Peekaboo-VL: Camera-Motion Text Hints and LoRA Letter Scoring for Drone Video Question Answering

Le Duc Minh
minh.leduc.0210@gmail.com
MAVU Workshop @ WACV 2027 — ReVA Drone Video Understanding Challenge

***Abstract*—Drone video is hard for video language models: the camera never stops moving, objects are a few pixels wide, and many questions ask *when* something happens. The organizers' model, ReMoSense, adds learned camera-motion tokens and reports a 3.9-point gain. We take the opposite route: instead of teaching the model motion, we *tell* it. A CPU probe estimates the frame-to-frame camera motion in milliseconds and renders it as one sentence in the benchmark's own answer vocabulary ("moving forward", "panning left"). Fine-tuning Qwen3.5-9B with LoRA at 32 frames and reading the answer from the letter scores A–D, our system reaches 83.6% on the ReVA test leaderboard, 3.5 points above ReMoSense, with no new network parts and one GPU. Most of that margin comes from the newer backbone and from fine-tuning on the test's videos, which already appear in training; the hint itself is cheap and easy to inspect, but its measured effect after fine-tuning is small. Beyond the headline number we report three findings that generalize. (i) The test set asks new questions about videos already in training, and the official validation set mostly repeats training; we build a development split with no training-answer leak that tracks the test score to about one point (a fifth of it are test questions that the validation set happens to label). (ii) A text-only model that never sees a frame scores 55.9%, so the real floor is 56%, not 25%. (iii) The wording of a text hint matters more than the measurement behind it: the same numbers in image terms cost a small model 12.5 points on camera-motion questions, while in the answers' own words they score 16 points higher than that. Temporal grounding remains the hardest task; its errors are near-misses of about one second, the gap between the timestamps the model sees.**

Code: https://github.com/LakoreAI/peekaboo-vl. Weights: https://huggingface.co/LakoreAI/peekaboo-vl.

***Index Terms*—video question answering, drones, UAV, vision-language models, camera motion, multiple-choice bias, LoRA fine-tuning**

I. Introduction

Drone video is an unfriendly input for a video language model. The camera never stops moving, often faster than anything on the ground, so "the car moves" and "the camera moves" look alike. The cars are a few pixels wide. And many questions ask *when* something happens, which needs a finer sense of time than the model usually gets.

The ReVA benchmark [1] tests exactly this, with eleven question tasks over clips from VisDrone, UAVDT, ERA and a new collection, Hawk-UAV. The organizers' own model, ReMoSense, extends Qwen2.5-VL-7B [2] with learned camera-motion tokens and object tokens and reaches 80.04% on test. Their ablation says the camera-motion tokens alone are worth 3.9 points overall and 9 points on temporal questions. Camera motion clearly matters, so the real choice is how to get it into the model.

This work was done by one independent researcher with a single rented GPU, so training new motion modules was not an option. We asked a simpler question: if the model cannot see the camera motion well, why not just tell it? Drone footage makes that cheap. The ground is far away and close to flat, so how one frame lines up with the next is described by a shift, a zoom and a rotation. Classic computer vision measures those in a few milliseconds on a CPU. We write the result as one line of text and let fine-tuning decide how much to trust it (Fig. 1). The same idea covers two other things the model cannot easily see: the true time of each frame, which recent Qwen models read as text timestamps [3], and the habits of whoever wrote the answers, which a tiny text-only model picks up surprisingly well.

We entered the ReVA challenge with one A40 GPU and a few weeks. This paper reports what that budget bought, including the experiments that failed. Our contributions are:

- **Tell, don't learn.** We inject camera ego-motion as text instead of adding learned motion tokens. A CPU probe measures the motion between sampled frames in a few milliseconds, and we render it in the benchmark's own answer vocabulary ("moving forward", "panning left"), so the model can match it to the options. Fine-tuned with LoRA on Qwen3.5-9B, the system reaches 83.6% on the ReVA test leaderboard, 3.5 points above the organizers' purpose-built ReMoSense-7B, using no new network parts and one GPU. We are explicit about where this margin comes from: the 9B backbone alone already matches ReMoSense, and fine-tuning on the seen videos adds the rest. After fine-tuning, the model does read the hint, but removing it costs only about 0.2 points overall.
- **A benchmark audit.** The ReVA test set asks new questions about videos that are already in training, and the official validation set is largely a copy of training questions. We build a development split with no training-answer leak that mirrors the test regime and tracks the test score to about one point. A fifth of it are test questions whose answers the validation set happens to contain, and we say so. We argue that ReVA results should be reported

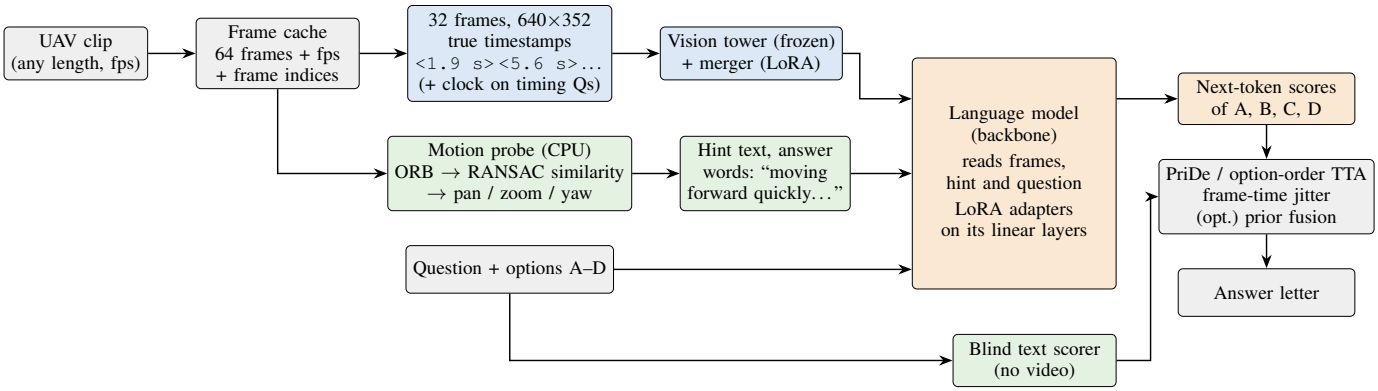


Fig. 1. The Peekaboo-VL pipeline. Each clip is decoded once into a cache that keeps its frame rate and which frames were taken, so the prompt can carry true times. A CPU step measures how the camera moves between frames and writes it as one line of text in the answers' own words. The backbone language model, fine-tuned with LoRA, scores the four letters. The final step can remove letter bias and, optionally, mix in a text-only prior.

alongside a trained blind baseline and per-task numbers.

- **Wording over measurement.** The same motion measurements, written in image terms ("the image shifts up") rather than the answers' words, *lowered* a small model's accuracy on camera-motion questions by 12.5 points. Written in the answers' words, the same numbers scored 16 points higher than that, and 3.4 above no hint at all. A text hint's wording is first-order, not a detail.
- **An empirical map, negative results included.** We measure what each part is worth with pre-registered success rules: the motion hint after fine-tuning, question packing, letter-prior debiasing and prior fusion, an object-motion hint, a caption hint from the benchmark's own question writer, a second training epoch, and frame density at test time. Temporal grounding stays the hardest task, and its errors are near-misses of about one second, the gap between the timestamps the model sees.

The rest of the paper follows the system: data and split (Sec. III), what it is built on and how it compares to prior work (Sec. II, Sec. IV), the method step by step (Sec. V), the main results (Sec. VI), and an ablation of every part (Sec. VII).

II. Related Work

Video language models. Qwen2.5-VL [2] tells the model the time of each frame through the way it numbers token positions (a "rotary position" scheme). Qwen3-VL [3] replaces this with plain text timestamps, such as "<3.0 seconds>", placed before each pair of frames. Qwen3.5 [4] reuses the Qwen3-VL vision part on a new language model. In that model, three out of every four layers are "Gated DeltaNet" [5] layers, which work like a running memory and are cheaper than standard attention; every fourth layer is standard attention. We build on Qwen3.5 and use its text timestamps as one of the things we "tell" the model.

Question answering on drone and remote-sensing video. ReVA [1] introduces the benchmark and the ReMoSense model. ReMoSense learns camera-motion tokens from cost volumes (a dense map of how well each patch matches the next frame) and object tokens from a recurrent transformer. Our motion hint carries the same kind of information, but as text, with no new model parts. EgoSIS [6] is the closest drone work we know of. It fits a robust image-plane transform between frames from optical flow, which is almost the same measurement our probe makes, and turns it into four learned "spatial evidence" tokens per slice of video on top of Qwen3-VL. On SIS-Bench this lifts an 8B model from 63.1% to 76.2%. So the two routes start from the same signal and differ only in how it reaches the language model: as learned tokens, or as a sentence. UrbanVideo-Bench [7] is another drone-video benchmark. Its videos are filmed from a drone flying through city streets, and its questions test whether a model can follow the flight, remember what it saw, and plan. Its 5,355 multiple-choice questions have between two and seven options. Unlike ReVA, it asks mainly about the drone's own movement and route, which is close to our camera-motion questions. We use it only as an outside test of whether our motion hint carries over to other drone videos; we never train on it.

Camera motion in video models. CameraBench [8] shows that classic geometry methods ("structure from motion") fail on camera movements that need understanding of the scene, while video models fail on simple geometric ones. Combining the two works better than either. Feng et al. [9] look inside a Qwen2.5-VL vision encoder and find that camera motion is only weakly represented, especially in its later layers. They extract camera motion with a 3D geometry model, turn it into motion labels with a small classifier, and give these labels to the video model in the prompt. This improves camera-motion answers without fine-tuning the video model. Their approach is the closest to ours: both *tell* the model the camera motion in text. Ours differs in three ways. We use a much simpler and faster estimator, because drone scenes are nearly flat. We write the hint in the exact words of the benchmark's answers. And we fine-tune the model together with the hint, so it learns when to trust it. Two recent papers learn the motion instead. CameraVLM [10] trains camera and optical-flow tokens together with the language model and argues that the motion has to be read by the language model, not estimated on its own. CamDistill [11] distills the camera tokens of a

3D model into the video model during training, so the 3D model is not needed at test time. It also tries our kind of shortcut, writing the 3D model's per-frame pose as text in the prompt, without fine-tuning. That helps the model find *when* the motion changes but not *which* motion it is. Our results point at the wording as one reason. Given to a small model without fine-tuning, the same motion numbers dropped camera-motion accuracy from 68.2% to 55.7% when written in image terms, and brought it to 71.6% when written in the answers' words (Sec. VII-B). Our estimator itself uses ORB features [12], simple matching, and a RANSAC fit [13].

Biases in multiple-choice answers. Language models prefer some option letters regardless of the content, for example "A" over "D". PriDe [14] measures this preference by showing the same question with the options in rotated orders, and then divides it out, without needing any correct answers. Ok et al. [15] split this kind of bias into two parts: how often an answer flips when the options are reordered, and how strongly the answers lean toward some letters. Prior correction such as PriDe mostly fixes the lean, while averaging over option orders mostly reduces the flips. We use both: PriDe and option shuffling at test time, and random option order during training.

A related, older problem is the *language prior*: patterns in the questions and answers that reveal the answer without the image. Goyal et al. [16] showed how strong such priors were in image question answering, and Chakraborty et al. [17] show that video models still lean on the wording of text options: one strong model drops from about 74% to about 30% when the same options are shown as frames instead. For aerial video in particular, a recent meta-analysis of 67 papers [18] found that none of the aerial question-answering benchmarks reports a blind (text-only) score, so there is no floor to read their accuracies against. Our trained blind baseline (Sec. V-L) gives that floor for ReVA.

III. Task and Data

A. The benchmark

Each ReVA record has a video path, a question, four options (A–D) and, for training data, the correct letter. The questions fall into eleven tasks in four groups (Table I). The challenge test set has 4,000 questions over 1,014 clips. The score is accuracy, overall and per task. Fig. 2 shows one training clip with three of its questions, and Fig. 3 shows one question from each video source.

What the tasks ask, in plain words. *General Understanding* asks what the video is about or what is in it. *Object & Land Cover* asks what kinds of objects or ground (grass, water, buildings) appear, and sometimes how many. *Change Detection* asks how something changes during the clip, such as how a car moves or how the view changes. *Temporal Grounding* asks when something happens, as a time or a time range. *Trend & Pattern* asks about repeated or overall movement, such as the main direction of traffic. The three spatial tasks ask where things are relative to each other, how the area is laid out (for example, a grid of roads), and how the camera looks at or moves over the scene. The three causal tasks ask why something happens, what will happen next, and what would happen if something were different. These last three rely mostly on general knowledge and are already easy for large models.

TABLE I
The eleven tasks and how many test questions each has.

Group	Task (leaderboard name)	Test Q
Factual Perception	General Understanding (FP_General)	180
	Object & Land Cover (FP_ObjectLand)	660
	Change Detection (FP_Change)	500
Temporal	Temporal Grounding (TU_Temporal)	640
	Trend & Pattern (TU_TrendPattern)	360
Spatial	Geometric Relation (SR_Geometric)	400
	Structural Layout (SR_Structural)	340
	Perspective & Viewpoint (SR_Viewpoint)	480
Causal	Causation (CR_Causation)	180
	Consequence (CR_Consequence)	100
	Hypothetical (CR_Hypothetical)	160

TABLE II
What the released annotation files show.

Statistic	Value
Train / val / test questions	15,773 / 2,000 / 4,000
Train / val / test clips	1,045 / 379 / 1,014
Test clips also in train	1,012 / 1,014
Test questions on train clips	3,997 / 4,000
Test exact copies of train questions	24
Test close copies (word overlap ≥ 0.6 / ≥ 0.8)	334 / 67
Val exact copies of train questions	1,572 / 2,000
Correct answer is the longest option (train+val)	36.3%
Our fit / dev split (questions)	14,068 / 2,078
Dev questions on fit clips	2,074 / 2,078

B. What the data files reveal: same videos, new questions

Before building anything, we checked the released files carefully. Table II summarizes what we found.

The most important fact is this: 1,012 of the 1,014 test clips, and all 529 longer source videos they were cut from, also appear in the training set. 3,997 of the 4,000 test questions are about a clip that has training questions. But the test questions themselves are new. Only 24 are exact copies of a training question on the same clip, and 334 are close copies (at least 60% of their words in common). So the test asks *new questions about videos the model has already seen in training*.

The official validation set is different. 1,572 of its 2,000 questions repeat a training question on the same clip, and 1,569 of those have the same answer. A model trained on the training set will look good on this validation set simply because it remembers the answers. That number would tell us little about the test.

C. A development split that behaves like the test

Because the official validation set cannot be trusted, we built our own split. Our *dev* set has two parts: the 428 validation questions that are *not* copies of training questions, and 150



Task	Question	Options (correct one in bold)
TU_Temporal	At what point in the video do the white tents first appear in the frame?	A. 4.8 s **B. 3.5 s** C. 2.0 s D. 0.0 s
SR_Viewpoint	What is the primary camera motion state throughout the video?	A. Panning left/right **B. Moving forward** C. Zooming in/out D. Rotating
FP_Change	How does the camera's perspective change throughout the video?	**D. From a wide aerial view of the neighborhood to a closer view of the field and tents** (3 other options)

Our motion hint for this clip: "Camera motion: 0.1–5.9s moving forward quickly while rotating clockwise. Mostly moving forward (100% of the clip)."

Fig. 2. A ReVA training clip (`Hawk_UAV/PH/DJI_0229_d4_01`): six evenly spaced frames with their true times, three of its questions, and the motion hint our system writes for it. The tents enter the frame at t=3.5 s, which is the answer to the temporal question. The hint already contains the answer to the camera-motion question, in the option's own words.



ERA, CR_Causation. Likely cause of the many parked cars? **(B) Baseball games** (vs. restaurant, highway, mall).



VisDrone, FP_ObjectLand. What type of land cover is visible on the left side of the parking lot? (A) A paved road (B) A body of water (C) A dense forest **(D) A field of green grass**.



UAVDT, TU_Temporal (23.2 s clip). At what time does the white van first appear in the video? (A) 0:08 (B) 0:12 **(C) 0:03** (D) 0:05.



Hawk-UAV, SR_Geometric. How is the tennis court positioned relative to the soccer field? (A) Above (B) To the left (C) Below **(D) To the right**.

Fig. 3. One example from each of the four video sources, three evenly spaced frames each with their true times (correct answers in bold). ERA clips are square and about 6 s long; UAVDT clips can last tens of seconds; temporal questions use both seconds and `m:ss` formats.

randomly chosen training questions per task, set aside. Our *fit* set (the data we train on) is everything else, after removing every record that asks a dev question about the same clip. The clips are still shared between fit and dev, just as they are shared between train and test. This copies the "same video, new question" setting of the real test.

One property of this split must be stated plainly. The 428 validation questions that are not copies of training questions turn out to be copies of *test* questions: all 428 have the same clip and the same question text as a test question. So 20.6% of dev are test questions whose answers the official validation file contains. They are not training answers, so dev has no training-answer leak, but dev is not disjoint from the test, and its agreement with the test score is partly built in. The other 1,650 dev questions are held-out training questions.

The dev set has 2,078 questions. With that many questions, the accuracy has a margin of error of about $\pm$0.9 points. Dev tracks the test well: our blind scorer gets 55.5% on dev and 55.9% on the test leaderboard, and the test-weighted dev accuracy (below) lands within 0.2 points of the test for our first fine-tuned model and within about 1.2 points for the final one (Sec. VI).

One more detail matters. The tasks do not have the same share of questions in dev and in test. Dev has 150–231 questions per task, while the test has 100 to 660. Some tasks that are common in the test, such as Temporal Grounding (640 questions), are among the hardest. So we also compute a *test-weighted* dev accuracy, where each task's dev accuracy is weighted by its number of test questions. This is the number we use to pick between systems.

D. The clips are not all alike

The dataset card describes 15-second clips at 640×360 pixels. Decoding all 1,047 clips used by the questions shows a different picture (Table III). ERA clips are square and under 6 s. Hawk-UAV clips range from 0.1 to 40 s. Ten clips are shorter than one second, yet they carry 137 questions between them; those questions cannot be answered from the video. Because we always show the model the same number of frames (16 in our first runs, 32 in the final model), the time between two frames ranges from a fraction of a second to more than two seconds. This is why the model must be told the *true* time of each frame, computed from each clip's own frame rate.

E. Training data balance

The task mix of the training data differs from the test. Temporal Grounding and Object & Land Cover are 12.2%

TABLE III
CLIPS USED BY THE QUESTIONS, BY SOURCE (DECODED FROM THE RELEASED FILES).

Source	Clips	Res.	fps	Length (s) min/med/max
Hawk-UAV	556	640×360	30	0.1 / 8.0 / 40.2
ERA [19]	261	640×640	∼24	5.4 / 5.8 / 5.8
VisDrone [20]	183	640×360	30	0.2 / 6.0 / 6.0
UAVDT [21]	47	640×360	30	3.9 / 16.0 / 35.0

and 11.1% of our fit set, but 16.0% and 16.5% of the test. Both are weak tasks for our model. This suggests giving these tasks more weight during training, which is one of the next experiments.

IV. BACKGROUND

A video language model does not watch a whole clip. It reads a small number of frames, 16 or 32 in this paper, taken at even intervals. Each frame is cut into patches, a vision encoder turns the patches into vectors, and the vectors are placed into the language model's input next to the words of the question. Qwen3-VL and Qwen3.5 also write a short text timestamp (for example `<3.0 seconds>`) before each pair of frames, which is how the model knows *when* each frame was taken.

This design has two weak points on drone video. The motion between frames is never given directly: the model must infer it from differences between frames that may be a second apart. And with 16 frames the model sees roughly one timestamp per second on a typical clip (two per second at 32), so it cannot separate events closer together than that. The rest of the paper addresses both.

Our remaining design choices are standard, so we fix them here and use them throughout. We score the four answer letters instead of generating an answer (Sec. V). We fine-tune with LoRA [22]. We measure the video's true contribution against a blind text-only baseline (Sec. V-L), following the lesson of Goyal et al. [16] that question sets carry answer patterns which need no image. And we compare two systems with a paired exact McNemar test, under success rules written down before each run.

V. METHOD

Fig. 1 shows the whole pipeline. We explain it in the order in which a question goes through it.

A. Step 1: decode each clip once

Decoding video is the slowest part of reading it, and each clip is used by about 15 questions. So each clip is decoded once into a cache of 64 evenly spaced frames. For all 1,047 clips this takes under a minute and about 7 GB of disk; after that, loading a clip's frames is a plain file read. Together with the frames we store the clip's frame rate and the exact frame numbers we took. From these, the true time of every frame is known. The final model takes 32 of the 64 frames (our first runs used 16). Each frame comes from one of 32 equal time segments, a standard choice from action recognition [23]. A "phase" setting decides where inside its segment a frame is taken. At test time we take the centre. During training we draw a new random phase for every question, so the model sees many different sets of timestamps for the same clip. A model trained on one fixed set of frames gets worse when the frames shift (Sec. VII); random phases remove that weakness, and they act as a small data augmentation as well.

B. Step 2: two processor defaults that silently hurt

Two default settings in the Qwen3.5 video processor would have hurt us. Neither raises an error or a warning; we only found them by reading the prompt the processor actually builds.

Image size. The processor has a pixel budget (`size.longest_edge`, default 786,432 pixels). This budget covers *all* frames together, not each frame. With 32 frames, it would shrink each 640×360 frame to about 208×117 pixels, where a car is just a few pixels. We scale the budget with the number of frames, so each frame stays at 640×352. This gives 220 visual tokens for each pair of frames, about 3,500 for the 32 frames of a clip.

Timestamps. The processor computes frame times from video metadata. When frames are given without that metadata, it assumes 24 frames per second, so the times written in the prompt are wrong. Our cache passes the true frame rate and frame numbers. We checked the final prompt for a 25 fps, 15 s clip and it contains the correct `<1.9/5.6/9.3/13.1 seconds>` labels.

C. Step 3: measure the camera motion

For each pair of neighbouring sampled frames we do the following (Fig. 4).

1) Find a few hundred distinctive points (ORB key-points [12]) in both frames, such as road markings and roof corners.
2) Match points between the two frames. A "ratio test" throws away matches that are not clearly better than the second-best choice.
3) Fit one transform that explains most matches: a shift, a zoom and a rotation (a "similarity" transform with four numbers). RANSAC [13] handles the outliers: it tries many small random sets of matches and keeps the transform that agrees with the most of them. Matches on moving cars do not agree with the ground, so they are ignored.
4) Turn the four numbers into a label. The camera moves opposite to the scene: if the ground slides to the left, the camera is panning right. A camera tilted forward and flying forward sees the ground slide down and grow larger. Rotation means the camera turns. Small numbers mean the camera is still.

We then join neighbouring steps with the same label into timed segments, and fold segments shorter than 0.8 s into a neighbour. The whole probe takes a few milliseconds per frame pair on one CPU core; all 1,047 clips take under a minute.



Fig. 4. The motion probe between two sampled frames of the clip in Fig. 2. 304 of 450 point matches agree with one transform (green, 80 shown at random). They all stream downward, which is what forward flight looks like. Red dots are rejected matches, here on trees whose leaves all look alike.

Why not a learned motion model? Because this one needs no training, no GPU time and no labels, and its output is a sentence we can read and check. When it is wrong, we can see why (Sec. VI shows where it fails).

D. Step 4: write the motion in the answers' words

The first version of the hint described what happens *in the image*: "pans toward the top of the frame, fast and zooms in (descending)". Section VII-B shows that this wording confused the model. The final version uses the words of ReVA's camera-motion answers: *moving forward / backward*, *panning left / right*, *rotating*, *stationary*, with a speed word and a short summary at the end, for example:

```
Camera motion: 0.1–5.9s moving forward quickly
while rotating clockwise. Mostly moving forward
(100% of the clip).
```

E. Step 4b: a second hint, counts from a detector

The camera hint says how the *camera* moves, not how many objects are in the scene. Video models are known to undercount small, dense objects, and "Object & Land Cover" is one of ReVA's largest tasks. We therefore add a second hint produced by a VisDrone-trained detector (YOLO11x), the only learned component in the pipeline besides the model itself. The detector runs on the first, middle and last of the clip's cached frames, and its counts are merged into the words the questions use ("person", "car", "truck") and written as one sentence:

```
Detector counts: 0.0s car 23, pedestrian 4, truck
2; 3.0s car 19; 5.9s car 31, truck 3.
```

ReVA's counting answers are ranges ("10–20"), so a detector error of about 20% usually still lands in the correct range. The counts are computed once for all 1,047 clips and stored next to the motion hints, so training and testing add no detector cost. We built the hint for every clip but did not get to evaluate it within our compute budget, so the final model does not use it.

F. Step 4c: a clock drawn on the frames of timing questions

The timestamps in the prompt tell the model *when* each frame was taken, but only as text next to the frames. Temporal Grounding questions ask exactly this ("At what time does the red car first appear?"), and they are our weakest task. Following VTimeCoT [24], we also *draw* the time on the frames: a dark strip along the bottom edge, a red bar that fills from left to right as the clip plays, and the frame's time in seconds ("3.5s"). The model can then read the time from the image itself.

We draw the bar only for Temporal Grounding questions. Drawing it on every question helped timing questions but cost object questions, because the strip covers the bottom 16 pixels of the frame, where small cars and people often are. The bar is drawn the same way in training and in testing, so the model learns to use it. Tested on a model that was trained without it, the bar already raised Temporal Grounding by 0.9 points on dev.

G. Step 5: build the prompt

The model receives one user message: the 32 frames with their timestamps (with the drawn clock for Temporal Grounding questions), then the hint line, the question, the four options and a fixed instruction. Qwen3.5 can "think" in writing before answering; we switch this off, so the answer letter comes right after the start of the model's reply.

H. Step 6: score the four letters

At the answer position we read the model's scores (logits) for the tokens A, B, C and D and treat them as the scores of four classes. To save memory, we compute scores only for that one position. During training, the loss is the usual cross-entropy over these four scores. In other words, we train the model to put high probability on the correct letter, and we ignore all other words. There is no text generation and no answer parsing.

We chose this over generating an answer for three reasons. It never fails to produce a valid letter. It gives four scores per question, which the later steps (letter debiasing, mixing in a prior, combining models) need. And it is cheap: computing scores only for the last position avoids building the full 250k-word output layer for every token. Scoring takes 0.63 s per question on the A40 at 16 frames and about 1.4 s at 32 frames (Step 7b makes this faster).

I. Step 7: fine-tune with LoRA

We attach LoRA adapters [22] (rank r=64, scale α=128, dropout 0.05) to every linear layer of the language model. This covers the attention layers, the Gated DeltaNet layers and the feed-forward layers. We also attach them to the small "merger" that connects the vision part to the language part. The vision encoder itself stays frozen. In total, 161.6M of 9.57B parameters are trained (1.69%). The adapter file is about 620 MB, so keeping many experiments around is cheap, and training at 32 frames fits in 28 GB of the A40's 46 GB.

The final model trains on all 14,068 fit questions for one pass, with one question per sequence. A second pass did not help: with branching (Step 7b) and the learning rate decayed over two passes, the two-pass model ends at 82.96%

TABLE IV
TRAINING AND TEST SETTINGS OF THE FINAL MODEL.

Base model	Qwen3.5-9B, bf16, SDPA attention
Video input	32 frames, 640×352, true timestamps
Frame sampling	random phase per question (train), centre (test)
Hints	camera motion, in answer words; drawn clock on Temporal Grounding frames
Adapter	LoRA r=64, α=128, LM linears + merger
Data	all 14,068 fit questions, one per sequence
Optimizer	AdamW, lr 10^{-4}, 3% warm-up, cosine decay
Batch	1×16 accumulation, 880 steps (1 pass), clip 1.0
Training time	17 h (about 70 s per 16-question step)
Hardware	1× NVIDIA A40 (46 GB)
Speed	4.3 s per training question
Decoding	PriDe [14] (5% subset) or plain

test-weighted dev, against 83.04% for one pass. A faster recipe packs eight questions about one clip into one sequence (Sec. VII-E). Packing costs accuracy, and we now know why: each later question in the sequence can see the correct answers to the earlier ones. The model therefore trains on a prompt it never sees at test time, and with packing, more training data did not help. Step 7b gives the speed of packing without this leak.

Every time a training question is used, we shuffle its four options into a new random order. The correct answer is therefore equally likely to be A, B, C or D, and the model cannot learn to prefer a letter. Table IV lists all settings of the main run.

J. Step 7b: one video pass for many questions, without leaks

Each clip has about 15 questions, and the video makes up almost all of the prompt (about 3,500 of 3,700 tokens). Scoring the questions one by one therefore processes the same video 15 times. We split each prompt into a *prefix* shared by all questions about the clip (the video and the hint) and a short *suffix* per question (the question, the options and the start of the answer). The prefix runs through the model once, and every suffix then continues from it on its own branch:

- In the full-attention layers, each suffix attends to the prefix and to its own earlier tokens, never to another suffix.
- In the Gated DeltaNet layers, which keep a running state instead of attending, each suffix starts from the state the prefix ends in (and from the last few inputs of the layer's short convolution).

Every question therefore sees exactly the prompt it would see alone, so nothing leaks between questions, and during training the gradients of all branches add up in the shared prefix. We checked this against scoring the questions one at a time on Qwen3.5-0.8B in 32-bit floats: the letter scores agree to 10^{-4} and the adapter gradients to a relative 6×10^{-5}. In 16-bit floats the scores move by rounding only. Scoring eight questions per video pass was 3.7 times faster in our test, and the same split makes training about as fast as packing. Questions with and without the drawn clock see different frames, so they go into separate groups.

K. Step 7c: two more training signals (being tested)

Both signals change only the training loss; the test prompt stays the same.

The clip played backwards. Video models often give the same answer when the frames are shuffled or reversed, which means they answer from what is in the frames, not from their order [25]. For questions whose answer must change when time runs backwards, we also show the model the reversed clip and push it away from the original answer. These are Temporal Grounding questions about a moment ("when does the bus first appear"), Change Detection questions about a direction or a change, and camera-motion questions ("moving forward"). Questions whose answer stays the same, such as durations, a stationary camera or the viewing angle, are left out by simple text rules; about 16% of the fit questions qualify. The reversed clip keeps its timestamps and clock running forward, and it gets no motion hint, because the hint describes the original direction. With $p_{\text{rev}}(y)$ the probability of the original answer y on the reversed clip, the loss becomes

$$\mathcal{L} = \text{CE} + \lambda_{\text{rev}} \left(-\log\left(1 - p_{\text{rev}}(y)\right) \right), \qquad \lambda_{\text{rev}} = 0.5.$$

The extra term only says "not this answer"; it does not name a new one, and it fades once the original answer is unlikely. Before spending a run on it, we test whether the trained model already changes its answers on reversed clips; if it does, this signal has little to teach.

The clip's facts as plain text. The test set asks new questions about clips the model has already trained on. In language models, facts learned in a single format are often not recalled when the question is asked in a new way; training on the same facts in two formats at once fixes much of this [26]. Our training uses one format only: pick a letter for four given options. So each group of questions about a clip gets one extra branch (Step 7b) in which the model writes out that clip's training questions with their answers as plain text, without options or letters ("Q: When does the red car first appear? A: At 3.5 seconds"). The pairs are shuffled every time and capped at 24. The loss adds $\lambda_{\text{fact}} = 0.3$ times the mean word-level cross-entropy of this text. Because the text branch is a separate branch, the letter questions never see it; we checked that their scores are identical with and without it. Dev questions never appear in the text.

Outcome. The reversed-clip check ran first, on the final model's dev answers: on the 375 questions whose answer should change, the model kept its answer on the reversed clip 68.3% of the time, against 88.0% on 266 control questions whose answer should not change. The model already uses frame order where it matters, so by our pre-registered rule (a gap below 0.15 would justify the run; it is 0.197) we did not train with this signal. The fact-text signal was built and checked but not trained within our budget.

L. Step 8 (optional): remove letter bias and mix in a text prior

Letter bias. PriDe [14] estimates how much the model prefers each letter by scoring a small part of the test set with

the options rotated. It then divides this preference out of every score. It needs no correct answers.

Blind prior. We also train a small text-only scorer that never sees the video. For each option it computes 50 simple features: how long the option is compared with the others, how often similar answer texts were correct in training for this task, how many words it shares with the question, whether it is the largest or smallest number among the options, and a few task-specific copies of the strongest features. A logistic model over the four options learns how to weigh these features. It trains in seconds. At test time we can combine it with the video model as $\log p = \log p_{\text{VLM}} + \alpha \log p_{\text{blind}}$, where α controls how much we trust the blind scorer. We choose α on dev with cross-validation, with the folds split by clip so no clip is in both halves.

Plain decoding. Because the options are shuffled in training, the fine-tuned model has almost no letter preference left, so PriDe has little to correct. We therefore also tried plain scores: one pass with the official option order and no correction. On dev the plain scores looked slightly better, but on the test set PriDe won (83.55% vs. 83.25%). The plain answers lean slightly towards "A" (26.9% of answers), which is exactly what PriDe removes. We keep PriDe. A dev difference below about one point between two decodings of the same model is not a reliable guide.

Averaging over frame phases. Shifting all frames by half a segment changes about one answer in five on video benchmarks while leaving accuracy unchanged [27]. A model trained with random phases (Step 1) can be scored at several phases, for example 0.17, 0.5 and 0.83, and its answer probabilities averaged. This removes the dependence on one phase at the cost of one scoring pass per phase. We implemented it but did not evaluate it; the submitted system scores the centre phase only.

VI. Results

A. Experimental protocol

Unless stated otherwise, we report accuracy on our development split (Sec. III), which mirrors the challenge's "seen video, new question" test regime. Because dev's task mix differs from the test's, we also report a *test-weighted* accuracy that reweights the eleven tasks by their test question counts. It predicted the leaderboard within 0.2 points for our first model and within about 1.2 points for the final one, and it is our model-selection metric. Comparisons between two systems use a paired exact McNemar test over the questions the two answer differently. Each configuration is a single run at seed 42: we did not have the budget for repeated seeds, so our claims rest on paired tests and on decision rules written down before the run, not on seed averaging. All models are Qwen3.5-9B in bfloat16 with LoRA adapter weights on the language model and the vision–language merger (161.6M of 9.57B parameters, 1.69%); training runs on one NVIDIA A40 (46 GB). Decoding is either the official option order, the four cyclic option permutations averaged (4× the cost), or PriDe debiasing [14] (Sec. V-L). We report failed experiments alongside successful ones. Two systems appear below: our *first model* (16 frames, 8,000 fit questions), which was our first leaderboard entry and the reference for most ablations, and the final *Peekaboo-VL* model (32 frames with random training phases, the drawn clock on timing questions, all 14,068 fit questions; Table IV). Recall that 428 dev questions are test questions (Sec. III).

TABLE V
Overall accuracy (%). Dev unless marked; ReMoSense numbers are from its paper (test set).

System	Dev	Test
Random guessing	25.0	25.2
Longest option / answer-text prior	41.3 / 40.7	–
Blind text scorer (ours)	55.5	55.9
Qwen3.5-0.8B, no fine-tuning†	71.6	–
+ blind prior (α=0.75 fixed)†	77.6	–
Qwen3.5-4B, no fine-tuning	77.0	–
Qwen3.5-9B, no fine-tuning	79.6	–
+ blind prior (α by cross-validation)	80.1	–
ReMoSense-7B, fine-tuned [1]	–	80.0
First model (16 frames, 8k questions) + PriDe	84.7	82.6
+ blind prior (α by cross-validation, not used)	85.2	–
Peekaboo-VL (32 frames, clock, all fit), plain	**85.4**	83.3
+ PriDe letter debiasing (submitted)	85.0	**83.6**

† On a 250-question dev subset (25 per task, no Viewpoint).

B. How far do you get without the video?

Table V shows the main steps, from guessing to our final system. Without any video, the blind scorer reaches 55.5% on dev and 55.9% on the test leaderboard. That is more than double random guessing: over half of ReVA can be answered without looking at a single frame. Without any fine-tuning, Qwen3.5-9B already reaches 79.6% on dev, about the same as ReMoSense after its full fine-tuning (80.0% on test).

C. Which tasks are still hard?

Fig. 5 and Table VI compare the tasks.

The blind scorer does best where the same answers come up again and again: Viewpoint (72% on both dev and test) and General Understanding (71–72%). Camera-motion questions, for example, are mostly answered "moving forward". It does worst where the answer must be seen: Geometric Relation (37% dev, 44% test) and Change Detection (43% / 42%).

The 9B model without fine-tuning is close to perfect on the causal tasks (92–94%) and on General Understanding (97%). These questions mostly need common sense, which a large language model already has. Its weak spot is Temporal Grounding: 57.5%, twenty points below its average, followed by Change Detection (70.6%) and Trend & Pattern (72.8%). Temporal Grounding is also where ReMoSense makes the most errors on the test set (152 of its 798 errors).

D. What does fine-tuning add?

One pass of LoRA training over 8,000 fit questions (500 optimizer steps, 5.2 h on the A40) lifts Qwen3.5-9B from 79.6% to **84.7%** on dev (82.3% test-weighted). The model's

TABLE VI
ACCURACY PER TASK (%). DEV FOR OUR MODELS (2,078 QUESTIONS); TEST FOR THE BLIND SCORER'S LEADERBOARD ENTRY AND FOR REMOSENSE [1]. BEST PER COLUMN IN BOLD.

System	GU	OR	CD	TG	TP	GR	SL	PV	CA	CO	HR	Overall
Blind text scorer (test)	70.6	48.0	42.4	62.0	68.1	43.8	56.5	71.9	59.4	46.0	45.6	55.9
Qwen3.5-4B, no fine-tuning	94.8	80.2	67.1	51.1	73.9	68.5	80.0	71.2	90.8	91.3	90.5	77.0
Qwen3.5-9B, no fine-tuning	96.8	81.9	70.6	57.5	72.8	76.4	80.0	73.3	93.6	92.5	**91.7**	79.6
First model (9B, fine-tuned, 16 frames)	**99.4**	**83.9**	**80.1**	67.6	**82.6**	80.9	88.2	83.8	94.2	92.5	88.2	84.7
Peekaboo-VL (final, 32 frames, PriDe)	98.1	81.5	79.2	70.3	81.0	**84.3**	**88.8**	**85.3**	91.9	94.4	90.5	**85.0**
ReMoSense-7B, fine-tuned (test)	96.1	78.0	70.9	**76.2**	76.4	74.3	83.3	83.7	**95.3**	**97.7**	90.7	80.0

GU General, OR ObjectLand, CD Change, TG Temporal Grounding, TP Trend & Pattern, GR Geometric, SL Structural, PV Viewpoint, CA Causation, CO Consequence, HR Hypothetical.

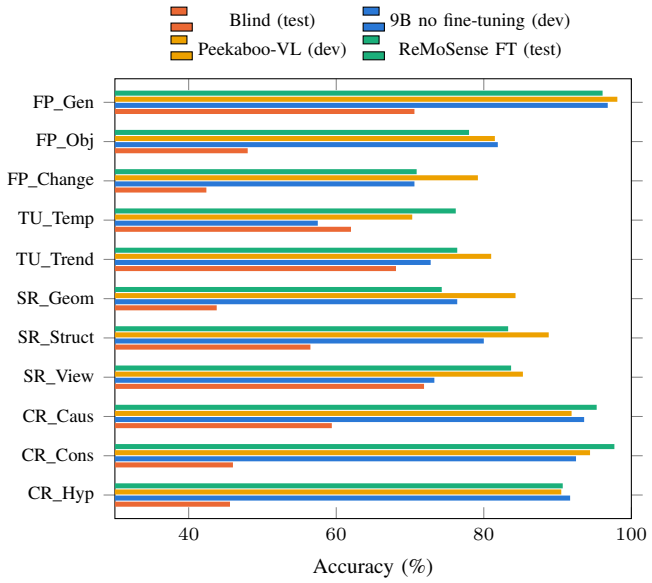


Fig. 5. Accuracy per task: our blind scorer on the test leaderboard, Qwen3.5-9B before and after our fine-tuning (final Peekaboo-VL, PriDe) on dev, and ReMoSense after fine-tuning on test. Exact values are in Table VI.

predicted letters are spread evenly over A–D (the KL divergence to a uniform spread is 0.0007, essentially zero). Because the options were shuffled every time a question was used, the trained model has no letter preference left to remove: PriDe leaves accuracy unchanged (84.74% both ways). On the test set this model scores 82.6%. Raw dev accuracy is 2.2 points higher than that, because dev has relatively more easy questions; the test-weighted dev accuracy (82.4% with PriDe) lands within 0.2 points of the test score, which is why we choose systems by it.

The final model changes three things: 32 frames with a random phase per training question (Step 1), the drawn clock on Temporal Grounding frames (Step 4c), and all 14,068 fit questions instead of 8,000. It reaches 85.4% on dev with plain scores (83.0% test-weighted) and **83.6%** on the test with PriDe, one point above the first model. Here the test-weighted dev under-predicted the test by about one point, so we read dev differences below one point as noise. Temporal Grounding rises from 67.6% to 70.3%, Geometric Relation from 80.9% to 84.3%, and Viewpoint from 83.8% to 85.3% (Table VI).

Table VI shows where the gains come from. They are largest exactly where the model without fine-tuning was weak: Viewpoint +10.5, Temporal Grounding +10.1, Trend & Pattern +9.8, Change Detection +9.5 and Structural Layout +8.2. Compared with ReMoSense's per-task test numbers, the dev accuracy of both our models is higher on seven of the eleven tasks. The clear exception is Temporal Grounding (70.3 for the final model vs. 76.2), which we study in Section VI-E. The causal tasks are within a few points and hold few errors.

Why does so little training help so much? Recall that the test asks new questions about videos that are in the training set. Each training question teaches the adapter a few facts about one video: where the river is, what colour the bus is, when the van appears. New questions about the same video can then reuse these facts. This is allowed by the challenge rules, and the organizers confirmed that training on these labels is permitted. What is *not* allowed, and what we do not do, is to look up training answers about a test clip at test time ("video memory").

This also tells the reader how to read our margin over ReMoSense. The 9B backbone without fine-tuning already matches ReMoSense, and fine-tuning on the seen videos adds about five points. The motion hint's share is small: removing it from the fine-tuned model costs about 4 of 2,078 dev questions (Sec. VII-B). The margin is mostly a newer backbone plus this training, not evidence that text hints beat learned motion tokens.

E. Why is temporal grounding still hard?

To understand the weak Temporal Grounding score, we split the 9B model's answers before fine-tuning by the type of answer. Questions with an exact time ("at 3.5 seconds", 102 dev questions) reach 53.9%. Questions with a time range ("from 1.0 s to 4.0 s", 57) reach 50.9%. Questions with a relative answer ("around the middle", 60) reach 70.0%.

When the model picks a wrong exact time, it is usually close: the median error is 1.0 s, and 62% of wrong answers are within 1 s of the correct time. That is about the same as the gap between the timestamps the model sees. With 16 frames there is one timestamp per pair of frames, so every 0.75–1 s on a typical clip. The model simply cannot tell 3.0 s from 3.5 s.

Section VII-C tests whether giving the model more time detail at test time helps.

F. Where are the remaining errors?

The fine-tuned model makes 317 errors on dev. Temporal Grounding alone has 70 (22%), followed by Change Detection (45), Object & Land Cover (39), Geometric Relation (36), Viewpoint (31) and Trend & Pattern (30). The three causal tasks together have 42, and General Understanding only 3. Fig. 6 shows three typical failures.

(a) *A broken clip.* The clip is cut to 0.23 s (7 frames), but the correct answer says a van is visible for 4 s. No model can answer this. The ten clips shorter than one second carry 137 such questions.

(b) *Counting at dusk.* The vehicles are small and partly hidden, and it is unclear which ones count as "parked". The model counts too few. This is a known weakness of video models, and a hint with counts from an object detector is one way to address it.

(c) *Where exactly is it?* The model gets the bus's motion right (moving away) but puts it at the side of the road instead of the centre lane. This is a viewpoint error, and a camera-motion hint does not help with it.

G. How good is the motion measurement on its own?

Separately from any video model, we checked the motion probe against the correct answers. We used the 862 training and validation camera-motion questions whose answer is one of six classes: forward, backward, pan left, pan right, rotating or stationary. From the probe's summary numbers alone, we picked a class for each clip (Table VII).

The probe picks the correct class 73.0% of the time. Always saying the most common class ("forward") would give 57.0%. When we choose the option that matches the probe's class, it answers 66.5% of these questions correctly (chance is 25%). Forward flight is recognized in 454 of 491 cases.

Pans are the weak spot. In 60 of 111 "panning right" clips, the forward part of the motion is larger in the clip summary, so the probe says "forward". We tried twice to fix this, both times with a pre-registered plan, and both attempts failed or were inconclusive. Features that describe the path separate the classes on average (the ratio of sideways to vertical movement is 0.49 for pans vs. 0.14 for forward flight), but the two groups overlap too much. Looking at the videos showed a deeper problem: on tilted views, turning the camera (yaw) changes the perspective of the scene. Our four-number model cannot represent that, so it reports a mix of pan and rotation.

VII. Ablation Study

This section changes one part of the system at a time and measures what it is worth. Not every change could be tested on the full 9B model: the first hint experiments ran on the small 0.8B model, and the newer training experiments use the faster packed recipe (Sec. VII-E) instead of the 5.2-hour first 9B run, and the last block uses the final recipe. So each row in Table VIII names its own reference, and only rows with the same reference should be compared with each other.

TABLE VII
MOTION PROBE ALONE VS. CORRECT CAMERA-MOTION ANSWERS (862 TRAIN+VAL QUESTIONS; ROWS: CORRECT CLASS, COLUMNS: PROBE).

	back.	forw.	pan L	pan R	rot.	still
backward (n=23)	**13**	5	1	0	0	4
forward (n=491)	4	**454**	6	10	7	10
pan left (n=40)	4	15	**9**	2	4	6
pan right (n=111)	7	60	4	**25**	6	9
rotating (n=23)	1	5	2	1	**13**	1
stationary (n=174)	11	39	1	3	5	**115**

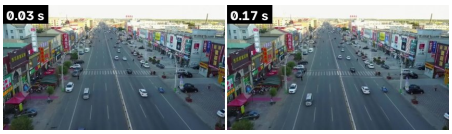
TABLE VIII
ALL ABLATIONS IN ONE PLACE. EACH BLOCK CHANGES ONE PART OF THE SYSTEM; "REF." IS THE REFERENCE SCORE (%) AND Δ THE CHANGE IN POINTS. FT = FINE-TUNED, TW = TEST-WEIGHTED DEV.

Change	Measured on	Ref.	Δ
Language model (Sec. VII-A); reference Qwen3.5-9B			
Cosmos-Reason2-8B	no FT, dev	79.6	-2.7
Qwen3.8-27B, 4-bit	no FT, dev	79.6	-2.1
Qwen3.5-4B	no FT, dev	79.6	-2.6
Cosmos-Reason2-8B	FT packed, TW	81.6	-0.7
Camera-motion hint (Sec. VII-B); camera-motion questions (88)			
Answer words, not image words	0.8B no FT	55.7	$+15.9$
Answer-worded hint vs. no hint	0.8B no FT	68.2	$+3.4$
Hint of a wrong clip	9B FT	85.2	-13.6
Hint removed	9B FT	85.2	-4.5
Time detail at test time (Sec. VII-C); exact/range TG (159)			
32 frames instead of 16	9B FT	67.3	-5.0
Jitter, 3 frame positions	9B FT	67.3	-4.4
Text prior mixed in (Sec. VII-D)			
α=0.75 fixed in advance	0.8B no FT, subset	71.6	$+6.0$
α by cross-validation	9B no FT, dev	79.6	$+0.6$
α by cross-validation	9B FT, dev	84.7	$+0.4$
Training recipe (Sec. VII-E)			
8 questions per video (packing)	9B FT, TW	82.4	-0.8
Learning rate 2×10^{-4} (was 10^{-4})	9B FT packed, TW	81.6	-3.3
Object-motion hint (Sec. VII-F)			
Object-motion hint added	9B FT packed, TW	81.6	-0.9
On the final recipe (Sec. VII-G)			
Second training pass	final FT, TW	83.0	-0.1
Caption hint, dev (PriDe)	final FT	85.0	$+0.4$
Caption hint, test questions in val	final FT, 392 Q	80.6	-1.8
64 frames at test (TG, TP, CD)	final FT, 634 Q	76.7	$+0.5$
QA-generator model, Qwen3-VL-30B-A3B	no FT, dev	79.6	$+0.4$

A. Which language model?

Our method does not depend on one particular model: the hint, the letter scoring and the LoRA recipe work with any video language model that reads frames with timestamps. We therefore tried several backbones (Table IX).

Without fine-tuning, Qwen3.5-9B is the strongest. The larger Qwen3.8-27B, run in 4-bit to fit on our GPU, scores lower (77.5%), and the smaller Qwen3.5-4B scores 77.0%. We also tried NVIDIA's Cosmos-Reason2-8B [28], a Qwen3-VL-8B model with extra training on real-world physical video such as traffic and robots. Without fine-tuning it scores 76.9%, 2.7 points below Qwen3.5-9B; it is weaker on Viewpoint (-11.5) and the reasoning tasks, but slightly better on Trend & Pattern



(a) TU_Temporal, 0.23 s clip. How long is the white van visible? Correct: *0.0–4.0 s*. Predicted: *0.0–1.5 s*.



(b) FP_ObjectLand. How many parked vehicles are visible? Correct: *5–6*. Predicted: *3–4*.



(c) FP_ObjectLand. Where is the red bus as it moves through the intersection? Correct: *centre of the frame, toward the top*. Predicted: *right side of the road, moving away*.

Fig. 6. Three failures of the fine-tuned model on dev: (a) a question that cannot be answered because the clip is cut short, (b) counting too few small objects at dusk, (c) placing an object in the wrong spot even though its motion is right. Labels give true times.

TABLE IX
LANGUAGE MODELS WE TRIED (DEV ACCURACY %; TW = TEST-WEIGHTED). "PACKED" RUNS PUT 8 QUESTIONS PER VIDEO IN ONE TRAINING SEQUENCE.

Model	Dev	TW
No fine-tuning		
Qwen3.5-4B	77.0	–
Qwen3.8-27B (4-bit)	77.5	–
Cosmos-Reason2-8B	76.9	73.3
Qwen3.5-9B	79.6	–
Qwen3-VL-30B-A3B (FP8; + caption hint)	**80.0**	76.8
Fine-tuned, packed recipe		
Cosmos-Reason2-8B	83.4	80.9
Qwen3.5-9B	**84.3**	**81.6**
Fine-tuned, unpacked		
Qwen3.5-9B, first model (16 frames)	84.7	82.4
Qwen3.5-9B, final (32 frames, plain)	**85.4**	**83.0**

(+2.7).

We then fine-tuned Cosmos-Reason2-8B with exactly the same recipe as a Qwen3.5-9B run (8 questions per video in one sequence; Sec. VII-E). Fine-tuning helps it even more than Qwen (+7.6 test-weighted), but it still ends 0.7 points below Qwen3.5-9B with the same recipe and 1.5 points below our first unpacked Qwen3.5-9B model. Almost all of the gap comes from one task, Change Detection (−8.7). On Viewpoint, Trend & Pattern, Object & Land Cover and Temporal Grounding it is 1–2 points *better*, which makes it a promising partner for combining two models. Qwen3.5-9B remains our main model and is the one behind our test score.

B. Does the wording of the hint matter?

We first tested the motion hint with the small Qwen3.5-0.8B model, without fine-tuning. We used all 191 dev Viewpoint questions, 88 of which ask about camera motion, and a control set of 100 causal and general questions that should not depend on motion (Table X).

The first wording, which described motion in image terms, made things *worse*. Accuracy on camera-motion questions fell from 68.2% to 55.7%. We looked at the 18 questions that this wording turned from right to wrong. In 12 of them the correct answer was "moving forward". But the hint never said "forward". It said "pans" in 16 of them and "zooms in" in 14. The model simply matched words from the hint to words in the options ("Panning from left to right", "Zooming in").

TABLE X
EFFECT OF THE HINT'S WORDING, QWEN3.5-0.8B WITHOUT FINE-TUNING, DEV (%): NO HINT, IMAGE WORDING, CAMERA (ANSWER) WORDING.

Question set (n)	No hint	Image	Camera
Camera-motion Viewpoint (88)	68.2	55.7	**71.6**
Viewpoint-angle (90)	41.1	45.6	44.4
All Viewpoint (191)	56.0	52.4	**59.2**
Control: causal + general (100)	89.0	88.0	87.0

TABLE XI
DOES THE FINE-TUNED 9B MODEL RELY ON THE HINT? (DEV, ACCURACY %).

Subset (n)	Hint	Removed	Wrong clip's
Camera-motion Viewpoint (88)	**85.2**	80.7	71.6
Change Detection (231)	80.5	80.5	80.5
Other Viewpoint (103)	82.5	83.5	83.5
All other tasks (1,656)	85.4	85.5	–

The geometry was right, but the words pointed to the wrong option.

We then wrote the *same* measurements in the answers' own words. This raised camera-motion accuracy to 71.6%. Compared with the image wording, 20 questions went from wrong to right and 6 from right to wrong (p=0.009). Compared with no hint at all, the gain is small and could be luck (+3.4 points, p=0.69). The control questions did not change (89 / 88 / 87%), so the hint does not distract the model on unrelated questions.

Does the trained model actually use the hint? After fine-tuning, we scored dev again in two altered ways: with the hint removed, and with the hint of a different, randomly chosen clip (so every clip gets a wrong hint; Table XI). The model uses the hint only where it is relevant. On camera-motion questions, a wrong hint costs 13.6 points (13 questions lost vs. 1 gained, p=0.002), so the model reads what the hint says. Removing the hint costs 4.5 points (p=0.34). So after training, the frames carry most of the information, and the hint adds a modest amount where it matters. Change Detection, other Viewpoint questions and all other tasks do not change (within 1 point). Our pre-registered rule required a 5-point drop when the hint is removed, so we call this result inconclusive, but with strong evidence that the model reads the hint's content.

TABLE XII
MORE TIME DETAIL AT TEST TIME, DEV TEMPORAL GROUNDING (FINE-TUNED 9B, TRAINED AT 16 FRAMES; ACCURACY %).

Condition	All (219)	Exact/range (159)	Relative (60)
16 frames, centred (as trained)	**68.0**	**67.3**	70.0
32 frames	66.2	62.3	**76.7**
Jitter, 3 positions	65.3	62.9	71.7

C. Does more time detail at test time help?

Section VI-E showed that wrong exact times are near-misses of about one timestamp gap. So we tried giving the trained model more time detail at test time, on its 219 dev temporal questions (Table XII). We tried 32 frames instead of 16. We also tried "jitter": averaging the answers from three views whose frames are taken at 25%, 50% and 75% of each time segment. Both made exact-time and range questions *worse*: they fell from 67.3% to 62.3% (32 frames) and to 62.9% (jitter, p=0.039). Only the relative questions improved, by up to 6.7 points.

Our explanation: the model was trained on 16 centred frames, so it learned to read time from that exact layout of timestamps. Moving every timestamp at test time breaks what it learned. The fix must therefore come during training: train with 32 frames, or with randomly shifted frames, so the model learns to read any layout. The final model does both (Step 1), and its Temporal Grounding accuracy rises from 67.6% to 70.3%. For that model, going further to 64 frames at test time no longer hurts, but it does not help timing questions either (Sec. VII-G).

D. Does the text prior help a big model?

Mixing the blind scorer into the small 0.8B model (no fine-tuning) added 7.7 points on a 220-question dev subset, with α chosen by cross-validation (64.1→71.8, p=0.0015). We repeated the test on 250 new questions with α=0.75 fixed in advance, and the gain held (71.6→77.6, p=0.008). On both sets, accuracy is flat for α between 0.5 and 1.5 (Fig. 7), so the exact value is not critical.

With the 9B model, the same mix adds only 0.6 points (79.6→80.1, p=0.33). The bigger model already knows most of these text patterns. After fine-tuning the gain is +0.4 (84.7→85.2, p=0.19), and the chosen weight falls to α=0.25 in every fold. Our pre-registered rule was to use the mix only if it adds at least 1 point with p<0.05. It does not, so the submitted system does not use it.

E. Training recipe: packing and learning rate

One video, many questions per training example. Each training example spends about 73% of its length on video tokens: roughly 1,920 of 2,000 tokens. Every question about the same clip repeats the same video tokens. ReVA has about 15 questions per clip, so most of the compute is spent reading the same frames again and again.

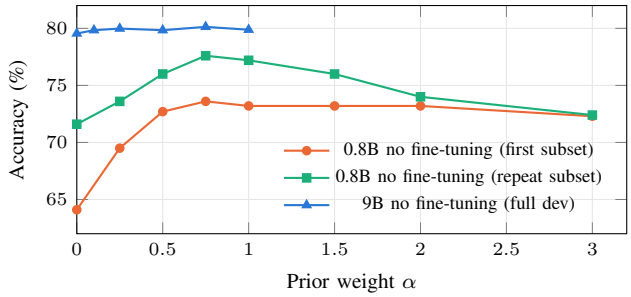


Fig. 7. Accuracy against the weight of the text prior. The gain is large and stable for the small model and almost disappears at 9B.

Question packing puts the video in once and then asks eight questions about it, one after another, as a normal chat:

$$\text{[video] hint } Q_1 \to A_1,\ Q_2 \to A_2,\ \ldots,\ Q_8 \to ?$$

The model is trained on the answer letter of every question. Because each question comes after the previous ones, no masking is needed, and the running memory of the DeltaNet layers is not a problem. We checked that the first question's scores are the same packed or unpacked (differences within bf16 rounding). Training becomes 6.4 times faster per question on a small model. The risk is a mismatch: during training, later questions see other questions and answers about the same video, which never happens at test time.

We repeated our main run with packing; the rule written beforehand was that accuracy must not drop by more than 1 point (test-weighted dev at least 81.4%). The packed run reaches 84.3% on dev and 81.6% test-weighted, 0.8 points below the unpacked run (82.4%) and inside the 1-point margin. Training took about 70 minutes instead of 5.2 h, 4.5 times faster. Temporal Grounding is slightly higher (68.5% vs. 67.6%). Packing passes, and because a run now takes about 70 minutes, the packed run is the reference for the training experiments below. The 0.8-point cost is most likely a leak: later questions in a pack can see earlier answers about the same video, which never happens at test time.

Learning rate. Every run so far used a learning rate of 10^{-4}, which we never tuned. A recent study that tuned ten LoRA variants found that, once the learning rate is tuned, plain LoRA is within one point of all of them [29]; the learning rate matters more than the variant. We therefore repeated the packed run with 2×10^{-4}, changing nothing else. It reached 78.3% test-weighted, 3.3 points below the packed run at 10^{-4} (81.6%), so we kept 10^{-4}. A lower rate (5×10^{-5}) was planned but not run.

F. Does an object-motion hint help?

The ReVA paper names one main failure of video models: confusing the camera's movement with the movement of objects. Our camera hint tells the model how the camera moves, but not how the cars move. The *object-motion hint* does this. For each pair of frames, we compute dense optical flow (how every pixel moves), fit the ground's movement with

a homography (a transform that also handles tilted views), and subtract it. Whatever still moves is a moving object. We keep only small moving blobs that continue in the same direction in the next step, and describe them by region of the frame and direction, in the questions' own words:

```
Object motion (camera motion removed): top-right:
moving right to left (0.9-3.9s).
```

This targets Change Detection and Trend & Pattern. The hint is computed for all 1,047 clips: 32 clips have no clearly moving objects and the median hint is 40 words long. We keep it only if it raises test-weighted dev by at least 0.5 points (with $p < 0.2$), or raises Change Detection plus Trend & Pattern by at least 2 points without any task losing more than 1.5.

Result: no gain. On the plain pass, test-weighted dev falls from 81.6% to 80.7% (-0.9). With letter debiasing on both runs it is unchanged (81.0 vs. 81.1); each run answers 62 questions the other misses (p=1.0). The tasks move in both directions: Structural Layout $+3.5$ and Geometric Relation $+2.8$, but Object & Land Cover -2.8 and Change Detection, one of the two target tasks, -1.3. Trend & Pattern gains only 1.1. Neither rule is met, so the hint is dropped. Two likely reasons, not yet checked: after fine-tuning, the frames already show the model how vehicles move, and on low, tilted views the residual flow still contains parallax from buildings that the blob filter does not fully remove.

G. What did not help the final model

The last block of Table VIII tests four more ideas on the final recipe. Besides dev, we use a second check that is closer to the test: the 392 test questions that match a validation question in clip, question and option set, so their answers are in the official validation file. The final model answers 316 of them correctly (80.6%).

A second training pass. Two passes with the learning rate decayed over both reach 82.96% test-weighted dev, against 83.04% for one pass, and 315 against 316 of the 392 test questions. One pass is enough.

A caption hint from the question writer. ReVA's questions were written by a large model from video captions. We replayed that captioning model (Qwen3-VL-30B-A3B) on every clip and gave its caption as a second hint. Without fine-tuning and without frames, the caption lifts the 9B model from 64.5% to 75.9% on dev (a different clip's caption: 59.3%), so the captions carry much of what the questions ask. After fine-tuning, however, the gain is gone: dev moves by $+0.4$ (85.5% vs. 85.0%, within noise), and on the 392 test questions the model with the caption is *worse* (309 vs. 316). We dropped it.

More frames at test time. Scoring the final model with 64 frames on its three weakest tasks leaves Change Detection unchanged (79.2%), moves Temporal Grounding from 70.3% to 69.9%, and raises Trend & Pattern from 81.0% to 83.2%. The last gain rests on 4 questions (p=0.125), so we do not count it.

The question writer as a second model. The same 30B model, without fine-tuning and with both hints, scores 80.0% on dev, below every fine-tuned 9B model. Combined with the final model (weights chosen on dev with folds split by clip), it adds 0.96 test-weighted points on dev (p=0.076), but only 1 of the 392 test questions. A combination with the caption-hint model looked better still on dev ($+0.84$, p=0.014) and answered one test question *fewer*. On this benchmark, gains of about one point on dev did not carry over to the test.

VIII. Discussion and Limitations

What the video adds. On ReVA, a text-only model already reaches 56% and a 9B video model without fine-tuning 80%. So most of the remaining room is in a few tasks: Temporal Grounding, Change Detection and Geometric Relation. These are exactly the tasks the blind scorer cannot solve. We suggest that results on ReVA be reported together with a trained blind baseline and per-task numbers, so readers can see what comes from the video.

Telling instead of learning. Text hints are cheap, easy to read and check, and work with any model. But they have two risks that learned motion tokens do not. First, their wording interacts with the answer options, as Section VII-B shows. Second, a wrong hint is stated just as confidently as a right one. Fine-tuning lets the model learn when to ignore the hint, and our test with wrong hints shows it does: it follows the hint on camera-motion questions and ignores it elsewhere.

Limitations. Most ablations ran on our first model (57% of the fit questions, 16 frames) or on the packed recipe, not on the final model, so their deltas are directions for the final recipe, not measurements of it. Our dev split shares clips with the fit set on purpose: it measures the "seen video, new question" setting of this challenge, not how well the model handles new videos. A fifth of dev are test questions whose answers the official validation file contains, which makes dev a closer proxy for the test than a disjoint split would be; we never used those answers to pick test predictions. We did not use the training answers about a test clip as context at test time ("video memory"); the organizers confirmed that this is not allowed, although training on those labels is. Finally, the probe's four-number motion model cannot represent camera turns on strongly tilted views; a homography, as used in our object-motion hint, is the natural next step. Our headline comparison is against ReMoSense's published numbers rather than a reproduction, and every configuration is a single seed; both are about our budget, and both are reasons to read the differences as directions rather than exact gaps.

Broader impact and ethics. This work studies aerial video understanding. Drone video is useful for disaster response, agriculture and traffic monitoring, and it also enables surveillance of people who did not consent to being recorded. The ReVA dataset is released for non-commercial research only and may not be redistributed, and we follow those terms and release no new video data. Our most reusable artifact is a few-millisecond CPU motion probe that runs on any video. The broader caution is methodological: when a benchmark's text alone is over half solvable and its test videos appear in training, reported accuracy can overstate genuine visual understanding.

A recent survey of aerial video benchmarks [18] reached the same conclusion from the other direction: none of the question-answering benchmarks it audited reports a blind floor.

IX. Conclusion

We set out to give a video language model the camera motion it cannot see well, by measuring it and writing it down instead of learning it. A few-millisecond CPU probe that describes the motion in the answers' own words, plus LoRA on an off-the-shelf video language model, scores 3.5 points above a purpose-built motion architecture on ReVA, with no new network parts and one GPU. Most of that margin comes from the newer backbone and from fine-tuning on the seen videos. The fine-tuned model does read the hint, but the hint itself adds little.

Three things we learned should hold beyond this benchmark. The test regime is seen-video, so a development split without training-answer leaks is essential. Text alone solves 56%, so the video's real contribution has to be measured against a blind baseline. And a hint's wording matters more than the measurement behind it. Temporal grounding remains the weakest task: training with 32 shifted frames and a drawn clock raised it by 2.7 points, but its errors are still near-misses of about one timestamp. We hope this encourages more video QA benchmarks to report text-only baselines and per-task numbers.

The code, the training recipe and the released LoRA adapter are available at https://github.com/LakoreAI/peekaboo-vl and https://huggingface.co/LakoreAI/peekaboo-vl.

References

[1] Z. Yao, L. Wang, Y. Yang, Z. Zheng, B. Lang, Q. Tang, J. Sheng, J. Xu, Y. Yang, J. Barker, X. Ying, and M. C. Chuah, "ReVA: A scene-centric dataset beyond repetition for remote sensing video question answering," 2026.

[2] S. Bai, K. Chen, X. Liu, J. Wang, W. Ge, S. Song, K. Dang, P. Wang, S. Wang, J. Tang *et al.*, "Qwen2.5-VL technical report," 2025.

[3] Qwen Team, "Qwen3-VL technical report," 2025.

[4] ——, "Qwen3.5: Native multimodal foundation models," https://huggingface.co/Qwen/Qwen3.5-9B, 2026, model card, accessed 2026-10-01.

[5] S. Yang, J. Kautz, and A. Hatamizadeh, "Gated delta networks: Improving Mamba2 with delta rule," in *International Conference on Learning Representations (ICLR)*, 2025.

[6] J. Yang, F. Ji, M. Cui, Y. Sun, H. Zhang, J. Zhu, and Y. Wang, "EgoSIS: From factorized visual ego-transitions to motion-canonical spatial evidence for UAV reasoning," 2026.

[7] B. Zhao, J. Fang, Z. Dai, Z. Wang, J. Zha, W. Zhang, C. Gao, Y. Wang, J. Cui, X. Chen, and Y. Li, "UrbanVideo-Bench: Benchmarking vision-language models on embodied intelligence with video data in urban spaces," 2025.

[8] Z. Lin, S. Cen, D. Jiang, J. Karhade, H. Wang, C. Mitra, T. Ling, Y. Huang, S. Liu, M. Chen, R. Zawar, X. Bai, Y. Du, C. Gan, and D. Ramanan, "Towards understanding camera motions in any video," 2025.

[9] H. Feng, S. H. Musunuri, and G.-M. Su, "Geometry-guided camera motion understanding in VideoLLMs," 2026.

[10] C. Cheng, R. Zhang, X. Li, A. Zhao, H. Wang, and X. Wang, "When left is right? camera-motion understanding beyond screen," 2026.

[11] D. Du, S. Du, J. Liu, Y. Yu, B. Gu, T. Han, H. Liu, E. Liu, Y. Zhang, X. Chen, and S. Guo, "Temporally grounded compositional camera motion understanding via geometric knowledge distillation," 2026.

[12] E. Rublee, V. Rabaud, K. Konolige, and G. Bradski, "ORB: An efficient alternative to SIFT or SURF," in *IEEE International Conference on Computer Vision (ICCV)*, 2011, pp. 2564–2571.

[13] M. A. Fischler and R. C. Bolles, "Random sample consensus: A paradigm for model fitting with applications to image analysis and automated cartography," *Communications of the ACM*, vol. 24, no. 6, pp. 381–395, 1981.

[14] C. Zheng, H. Zhou, F. Meng, J. Zhou, and M. Huang, "Large language models are not robust multiple choice selectors," in *International Conference on Learning Representations (ICLR)*, 2024.

[15] H. Ok, M. Park, and J. Lee, "One bias score, two factors: How mitigations affect option bias in multiple-choice evaluation," 2026.

[16] Y. Goyal, T. Khot, D. Summers-Stay, D. Batra, and D. Parikh, "Making the V in VQA matter: Elevating the role of image understanding in visual question answering," in *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2017, pp. 6904–6913.

[17] R. Chakraborty *et al.*, "Not another text benchmark: Putting the "visual" back in visual question answering for large video models," 2026.

[18] G. Angelidis, A. Psiris, G. Sfikas, M. Vakalopoulou, I. Varlamis, and G. T. Papadopoulos, "The gap is in the geometry: A convergent meta-analysis of multi-facet aerial and remote sensing video understanding," 2026.

[19] L. Mou, Y. Hua, P. Jin, and X. X. Zhu, "ERA: A data set and deep learning benchmark for event recognition in aerial videos," *IEEE Geoscience and Remote Sensing Magazine*, vol. 8, no. 4, pp. 125–133, 2020.

[20] P. Zhu, L. Wen, D. Du, X. Bian, H. Fan, Q. Hu, and H. Ling, "Detection and tracking meet drones challenge," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 44, no. 11, pp. 7380–7399, 2021.

[21] D. Du, Y. Qi, H. Yu, Y. Yang, K. Duan, G. Li, W. Zhang, Q. Huang, and Q. Tian, "The unmanned aerial vehicle benchmark: Object detection and tracking," in *European Conference on Computer Vision (ECCV)*, 2018, pp. 370–386.

[22] E. J. Hu, Y. Shen, P. Wallis, Z. Allen-Zhu, Y. Li, S. Wang, L. Wang, and W. Chen, "LoRA: Low-rank adaptation of large language models," in *International Conference on Learning Representations (ICLR)*, 2022.

[23] L. Wang, Y. Xiong, Z. Wang, Y. Qiao, D. Lin, X. Tang, and L. Van Gool, "Temporal segment networks: Towards good practices for deep action recognition," in *European Conference on Computer Vision (ECCV)*, 2016, pp. 20–36.

[24] J. Zhang, Y. Guo, R. A. Potamias, J. Deng, H. Xu, and C. Ma, "VTimeCoT: Thinking by drawing for video temporal grounding and reasoning," 2025.

[25] Z. Kang, S. Liu, T. Luo, W. Zhang, Y. He, L. Wei, G. Wang, L. He, S. Wang, G. Dong, J. Su, S. Chen, H. Ji, Q. Zhan, and K. Zhou, "Behavior pack optimization for video MLLM post-training," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2026, arXiv:2610.03141.

[26] Y. Zhang, B. Heinzerling, D. Li, and K. Inui, "Understanding fact recall in language models: Why two-stage training encourages memorization but mixed training teaches knowledge," 2025.

[27] L. Zhu, Y. Wang, Y. Lin, H. Li, and Y. Chen, "Stable scores, unstable answers: Frame phase and option order in video multiple-choice evaluation," 2026.

[28] NVIDIA, "Cosmos-Reason2-8B," https://huggingface.co/nvidia/Cosmos-Reason2-8B, 2025, model card; post-trained from Qwen3-VL-8B-Instruct; NVIDIA Open Model License.

[29] Lee, Ko, Chen, and Yeh, "Learning rate matters: Vanilla LoRA may suffice for LLM fine-tuning," *arXiv preprint arXiv:2602.04998*, 2026.